

Maximum Likelihood Analysis in R

Chad E. Brassil – June 29, 2007

Load special libraries

For the maximum likelihood analysis you will need download and load a library. I recommend `bbmle` which was developed by Ben Bolker.

```
library(bbmle)
```

Alternatively, you can use the older `stats4` library. Note `stats4` uses the function `mle()` whereas `bbmle` uses `mle2()`, and `stats4` does not have some of the supporting functions.

```
library(stats4)
```

A few special functions are included in the file `mleFunctions.R`. It can be loaded using the following code if a copy of the file is placed in your working directory.

```
source("mleFunctions.R")
```

Input data

First, use “File/Change Dir...” from the menu to set the correct working directory. Then you can load the data from a csv file. The top row should contain the names of each column. In this example the file is called `comboR.csv` and I store the data set into the variable `bristol`.

```
bristol <- read.csv("comboR.csv")
```

Optional

Look at the names of the columns

```
names(bristol)
```

Edit individual values

```
edit(bristol)
```

Attach the dataset so that you can refer to column name directly

```
attach(bristol)
```

Examine plots of columns

```
plot(bristol$clone,bristol$infection)
```

or, if dataset has been attached, then you can just use

```
plot(clone,infection)
```

Look at subset of data, in this case when species equals a or b

```
subdata <- subset(data,species == c("a","b") )
```

```
plot(subdata$clone, subdata$infection)
```

For an interactive 3-D examination of data try the feature-reach `scatter3D()`, which requires the `Rcmdr` library

```
scatter3d(data$mInfl,data$visits,data$height)
```

Create support function

To create the support function use the density form of a probability distribution (see the reference list of “Basic Distributions”). Setting `log = TRUE` will give in the log version and `-sum()` will take the negative sum across the data set. The resulting function should be the negative log likelihood.

```
ll <- function (m1,s)
  -sum(dnorm(bristol$infection,m1,s,log = TRUE))
```

Optional

The function `with()` can be used to temporarily attach a data set

```
ll <- function (m1,s) with(bristol,
  -sum(dnorm(infection,m1,s,log = TRUE)))
```

Categorical independent variables, i.e. factors

```
ll <- function (m.a,m.b,s) with(bristol,{
  m = c(m.a,m.b);
  -sum(dnorm(infection,mean=m[species], sd=s,
    log = TRUE))
})
```

However, be sure the vector, in this case `m`, is in the same order as

```
levels(bristol$species)
```

Interactions are done the same way, use “:” operator for interaction

```
with(bristol,levels(species:sex))
ll <- function (a.f,a.r,b.f,b.r,s) with(data,{
  m = c(a.f,a.r,b.f,b.r);
  -sum(dnorm(consumed,mean=m[sex:status],sd=s,
    log = TRUE))
})
```

Pick starting guess and find MLE

The Nelder-Mead method, or simplex method, is a generally robust method for finding the minimum of the support function.

```
guess <- list(m1 = 0.3, s = 0.1)
fit <- mle2(ll, start=guess, method="Nelder-Mead")
```

Optional

Parameter constraints can be given when using the L-BFGS-B minimization method

```
fit <- mle2(ll, start=guess, method="L-BFGS-B", lower
  = c(-1000,0), upper = c(1000,1000))
```

General summary of results

```
summary(fit)
```

Find confidence intervals for a model

```
confint(fit)
```

Or for a single parameter

```
confint(fit,"m1")
```

Examine likelihood profile plots. The profiles in display the transformation $|\sqrt{-L}|$ on the vertical axis. The second dashed, pink line from the top is the 95% confidence interval. The 0.99999 in conf forces it to plot a wider range of the profile; 0.95 plots the 95% confidence interval. Use par to show multiple plots at once.

```
par(mfrow=c(length(guess),1))
plot(profile(fit),conf=c(0.95,0.99999),absVal=TRUE)
par(mfrow=c(1,1))
```

To profile a single parameter

```
profile(fit,"m1")
```

To display the untransformed negative log likelihood, use the following function which is included in mleFunctions.R

```
prof <- profile(fit.saturating)
profpic(prof,"h",best=fit.saturating)
```

Visually examine model fit against data to ensure what you're doing makes sense

```
plot(mInfl,visits, xlab = "flowers", ylab = "visits")
```

Use curve() and coef() to show a line of your data onto of your data points

```
curve(coef(fit)["m"]*x+coef(fit)["b"], add=TRUE)
```

Use dev.copy2eps to save a copy of a figure

```
dev.copy2eps(file="fig.eps",width=8,height=6)
```

Randomly pick starting values and repeat optimization multiple times using this custom function called mle.random(). Output is the best fit. The support function is ll and the guess is the same as in mle(). Starting conditions are randomly selected n times from within *lower.guess* and *upper.guess*. Note, the code for mle.random() can be found in the file mleFunctions.R.

```
mle.random (ll, start, method, lower.start,
upper.start, n = 100 )
```

Example of use

```
ll <- function (m1,s)
-sum(dnorm(bristol$infection,m1,s,log = TRUE))
guess <- list(m1 = 0.3, s = 0.1)
fit <- mle2(ll, start=guess, method="Nelder-Mead")
summary(fit)
fit <- mle.random(ll,start=guess,lower.start=
c(-3,0),upper.start=c(3,10),n=100)
summary(fit)
```

Compare models

Extract AIC value for an individual model

```
AIC(fit)
```

To compare across a number of models, store each model in a separate fit value with an informative name. For example

```
AICtab(fit.constant,fit.linear,fit.saturating,
weights=TRUE,delta=TRUE,sort=TRUE)
```

To calculate AICc, use nobs to specify the number of observations. In the code below I extract that from data.

```
AICctab(fit.constant,fit.linear,fit.saturating,  
        nobs=dim(data)[1],weights=TRUE,delta=TRUE,sort=T)
```

The negative log-likelihood values for each function can be displayed using this code.

```
-logLik(fit.linear)
```

For nested models, you can use the Likelihood Ratio test. The user must ensure that models compared in this way are nested.

```
anova(fit.constant,fit.linear)
```